

# Introduction to Data Mining

## 06 - Clustering II

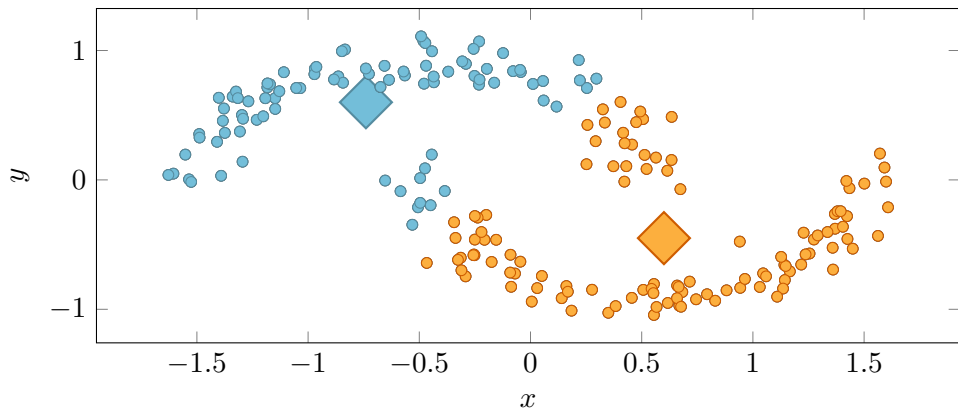
Benjamin Paaßen

WS 2023/2024, Bielefeld University

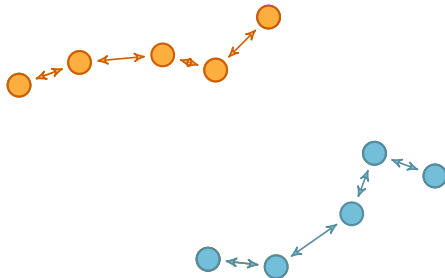
1. Agglomerative Clustering
2. Relational K-Means
3. Practical Story: Word Clustering

# Agglomerative Clustering

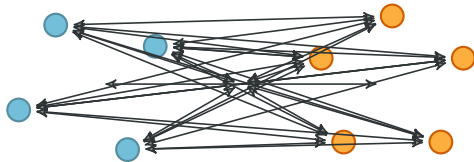
# Motivation: Non-spherical clusters



- Key idea: Each point is its own cluster; then merge closest clusters until only  $K$  clusters are left



# What does “closeness” mean?



- ▶ single linkage: distance between closest points; flexible cluster shapes, sensitive to 'bridges'
- ▶ complete linkage: distance between farthest points; spherical clusters
- ▶ centroid linkage: (squared) distance between means
- ▶ average linkage: all pairwise (squared) distances
- ▶ Ward's method: (squared) distance to shared mean (i.e. variance)

**function** AgglomerativeClustering(cluster distance function  $d$ , desired number of clusters  $K$ )

    Initialize  $\mathcal{C}_i = \{i\}$  for all  $i \in \{1, \dots, N\}$ .

**for**  $N - K$  repeats **do**

        Find  $k, l$  that minimize  $d(\mathcal{C}_k, \mathcal{C}_l)$ .

        Set  $\mathcal{C}_k \leftarrow \mathcal{C}_k \cup \mathcal{C}_l$ .

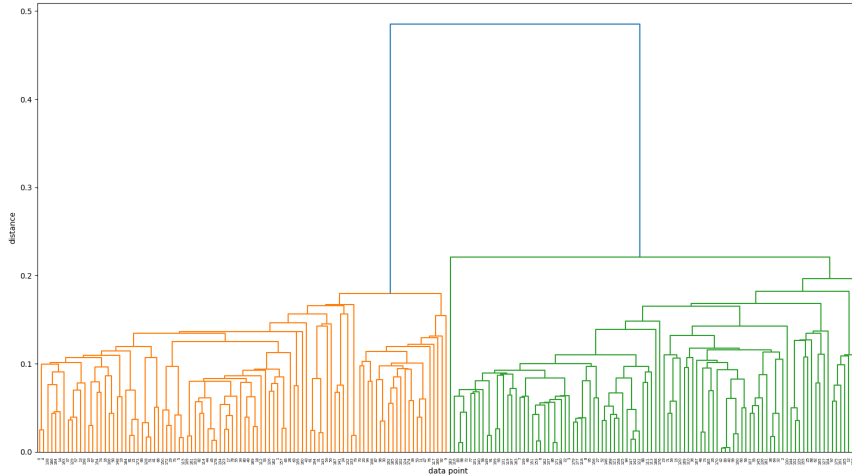
        Remove  $\mathcal{C}_l$  and update the cluster numbering.

**end for**

**return**  $\mathcal{C}_1, \dots, \mathcal{C}_K$ .

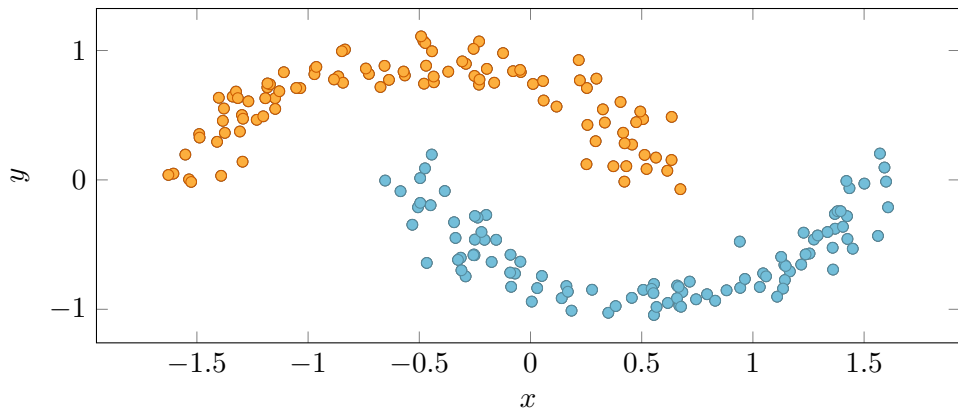
**end function**

# Single Linkage Dendrogram





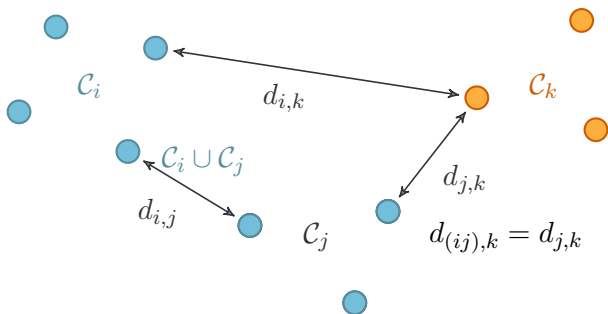
# Single Linkage Result



- ▶  $N - K$  iterations (one cluster vanishes per iteration)
  - ▶ computing  $d(\mathcal{C}_k, \mathcal{C}_l)$  takes  $\mathcal{O}(N^2)$
  - ▶ we need to evaluate  $\mathcal{O}(N^2)$  cluster-to-cluster distances
- ⇒ naive implementation can be as bad as  $\mathcal{O}(N^5)$  :(

## Example: Single-linkage clustering

- ▶ Key idea: Speed up computation by computing cluster-to-cluster purely based on prior cluster-to-cluster distances (recursively)



$$\begin{aligned}d_{(ij),k} &= \min\{d_{i,k}, d_{j,k}\} \\ &= \frac{1}{2} \cdot (d_{i,k} + d_{j,k}) - \frac{1}{2} \cdot |d_{i,k} - d_{j,k}|\end{aligned}$$

## Lance and Williams (1966)

Let  $d_{i,j}$ ,  $d_{i,k}$ , and  $d_{j,k}$  be the pairwise distances between clusters  $\mathcal{C}_i$ ,  $\mathcal{C}_j$ , and  $\mathcal{C}_k$ .  
Then, the distance  $d_{(ij),k}$  between  $\mathcal{C}_i \cup \mathcal{C}_j$  and  $\mathcal{C}_k$  is

$$d_{(ij),k} = \alpha_1 \cdot d_{i,k} + \alpha_2 \cdot d_{j,k} + \beta \cdot d_{i,j} + \gamma \cdot |d_{i,k} - d_{j,k}|$$

for suitable parameters  $\alpha_1, \alpha_2, \beta, \gamma \in \mathbb{R}$ .

variant	$\alpha_1$	$\alpha_2$	$\beta$	$\gamma$
single linkage	$\frac{1}{2}$	$\frac{1}{2}$	0	$-\frac{1}{2}$
complete linkage	$\frac{1}{2}$	$\frac{1}{2}$	0	$\frac{1}{2}$
centroid linkage	$\frac{N_i}{N_i + N_j}$	$\frac{N_j}{N_i + N_j}$	$-\frac{N_i \cdot N_j}{(N_i + N_j)^2}$	0
Ward's method	$\frac{N_i + N_k}{N_i + N_j + N_k}$	$\frac{N_j + N_k}{N_i + N_j + N_k}$	$-\frac{N_k}{N_i + N_j + N_k}$	0

- ▶ **note:** aggl. clustering requires no vectors, only distances!
- ▶ Dendrogram can be used to find no. of clusters
- ▶ With clever data structures, aggl. clustering can be improved to  $\mathcal{O}(N^2 \cdot \log(N))$ ; in practice, even faster (Bouguettaya et al. 2015)

## Relational K-Means

- Task: write a sorting program  $\Rightarrow$  How to cluster the answers?

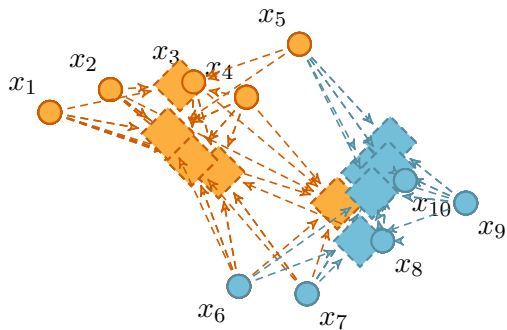
```
def bsort(A):  
    → for i in range(1, len(A)):  
        → for j in range(1, len(A)):  
            → if (A[j-1] > A[j]):  
                → tmp = A[j]  
                → A[j] = A[j-1]  
                → A[j-1] = tmp  
    → return A
```

```
def isort2(A):  
    → for i in range(1, len(A)):  
        → tmp = A[i]  
        → j = i-1  
        → while (j >= 0 and tmp < A[j]):  
            → A[j+1] = A[j]  
            → j -= 1  
        → A[j+1] = tmp  
    → return A
```

- We may not have a vector representation but pairwise distances (like tree edit distance)

- ▶ Basic idea: Same as  $K$ -means, but using only distances
- ▶ Trick 1: represent each prototype  $\mu_k$  only via coefficients  $\alpha_{k,1}, \dots, \alpha_{k,N}$ , such that
$$\mu_k = \sum_{i=1}^N \alpha_{k,i} \cdot x_i$$
- ▶ Trick 2: compute  $d(x_i, \mu_k)$  only based on distances between data points and  $\alpha_{k,1}, \dots, \alpha_{k,N}$





$\bar{i}$	$\alpha_{1,\bar{i}}$	$\alpha_{2,\bar{i}}$
1	0.05	0
2	0.05	0
3	0.05	0
4	0.05	0
5	0.05	0.05
6	0.05	0.02
7	0.05	0.02
8	0.05	0.05
9	0.01	0.05
10	0.01	0.05

- Can we compute a distance between a (weighted) mean and a point purely based on point-to-point distances?

## Relational Distance Formula (Hammer and Hasenfuss 2010)

Let  $x_1, \dots, x_N \in \mathcal{X}$ . Let  $\alpha_1, \dots, \alpha_N \in \mathbb{R}$  such that  $\sum_{i=1}^N \alpha_i = 1$ . Finally, let  $d : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}^+$  be a symmetric and self-equal function. Then, the point  $\mu = \sum_{i=1}^N \alpha_i \cdot x_i$  is well-defined and for any  $x \in \mathcal{X}$ , it holds:

$$d(\mu, x)^2 = \sum_{i=1}^N \alpha_i \cdot d(x_i, x)^2 - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \cdot \alpha_j \cdot d(x_i, x_j)^2 \quad (1)$$

- ▶ For full details, refer to Paaßen (2019, chapter 2.1)
- ▶ Starting point: for any symmetric and self-equal  $d$ , we can construct a bi-linear  $s : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ , such that  $d(x, y)^2 = s(x, x) - 2s(x, y) + s(y, y)$ .
- ▶ Starting from the right-hand-side, we obtain:

$$\begin{aligned} & \sum_{i=1}^N \alpha_i \cdot d(x_i, x)^2 - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \cdot \alpha_j \cdot d(x_i, x_j)^2 \\ &= \sum_{i=1}^N \alpha_i \cdot \left( s(x_i, x_i) - 2s(x_i, x) + s(x, x) \right) \\ & \quad - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \cdot \alpha_j \cdot \left( s(x_i, x_i) - 2s(x_i, x_j) + s(x_j, x_j) \right) \end{aligned}$$

$$\begin{aligned} &= \sum_{i=1}^N \alpha_i \cdot s(x_i, x_i) - 2 \sum_{i=1}^N \alpha_i \cdot s(x_i, x) + s(x, x) \cdot \left( \sum_{i=1}^N \alpha_i \right) \\ &\quad - \frac{1}{2} \sum_{i=1}^N \alpha_i \cdot s(x_i, x_i) \cdot \left( \sum_{j=1}^N \alpha_j \right) + \sum_{i=1}^N \sum_{j=1}^N \alpha_i \cdot \alpha_j \cdot s(x_i, x_j) - \frac{1}{2} \left( \sum_{i=1}^N \alpha_i \right) \cdot \sum_{j=1}^N \alpha_j \cdot s(x_j, x_j) \\ &= -2 \sum_{i=1}^N \alpha_i \cdot s(x_i, x) + s(x, x) + \sum_{i=1}^N \sum_{j=1}^N \alpha_i \cdot \alpha_j \cdot s(x_i, x_j) \\ &= -2s\left(\sum_{i=1}^N \alpha_i \cdot x_i, x\right) + s(x, x) + s\left(\sum_{i=1}^N \alpha_i \cdot x_i, \sum_{j=1}^N \alpha_j \cdot x_j\right) \\ &= -2s(\mu, x) + s(x, x) + s(\mu, \mu) = d(\mu, x)^2 \quad \square \end{aligned}$$

**function** Relational KMeans(matrix of squared distances  $D^2$  with  $N$  rows and columns, desired number of clusters  $K$ )

Randomly initialize  $A$  as  $N \times K$  matrix of positive numbers.

Divide columns of  $A$  by column sums.

**while**  $A$  still changes **do**

    Compute  $d(\mu_k, x_i)^2$  using Equation (1) for all  $i$  and  $k$ .

    Compute  $z_i \leftarrow \arg \min_k d(\mu_k, x_i)^2$ .

    Set  $\alpha_{k,i} = 1$  if  $z_i = k$  and 0, otherwise.

    Divide columns of  $A$  by column sums.

**end while**

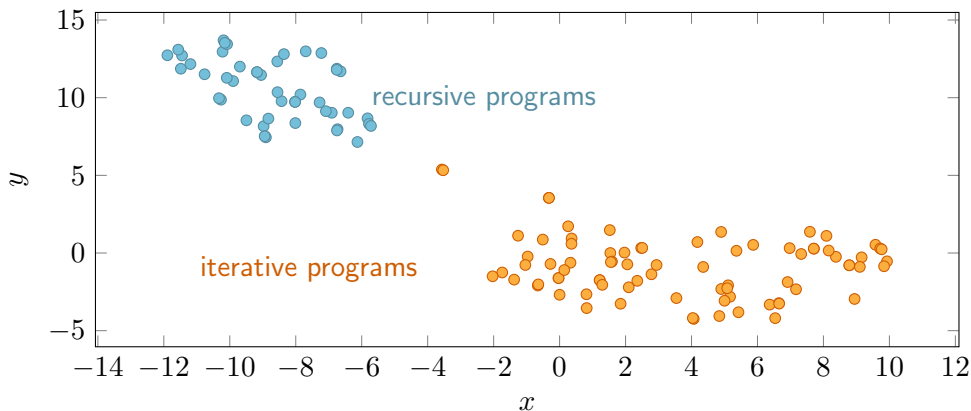
**return**  $A$ .

**end function**

- ▶ Almost the same as  $K$ -means, just with coefficient representation for prototypes and relational distance formula (1)
- ▶ Yields exactly (!) the same results as regular  $K$ -means for Euclidean distance

# Visualization: Programming Clustering

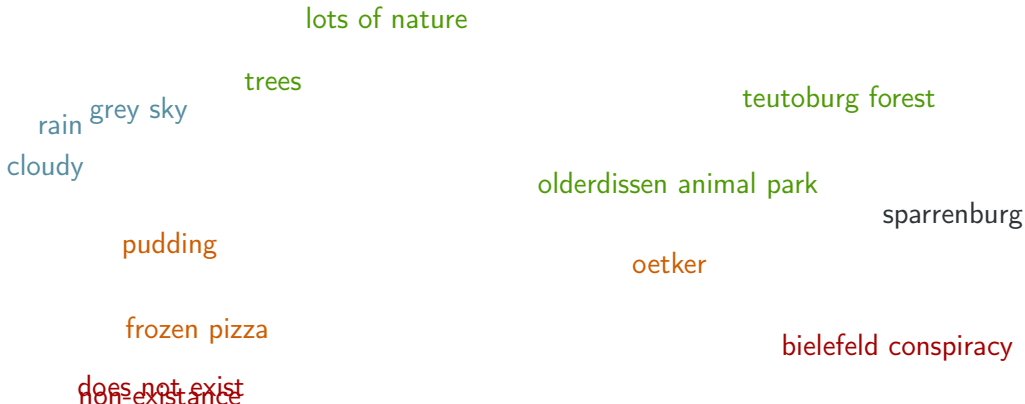
- ▶ distance function: adapted tree edit distance (Paaßen 2019)
- ▶ t-SNE embedding; color indicates relational  $K$ -means with  $K = 2$



## Practical Story: Word Clustering



- ▶ Without thinking too hard, what is the first word you associate with **Bielefeld**



- ▶ How to cluster these answers – without manual labor?

# Step 1: Define a distance function

Try to think of a good distance function for words

- ▶ bag-of-words?
  - ▶ edit distance/Levenshtein distance?
  - ▶ normalized compression distance?
- ⇒ cosine distance on **word embeddings** by BAAI/bge-large-en-v1.5 language model from huggingface

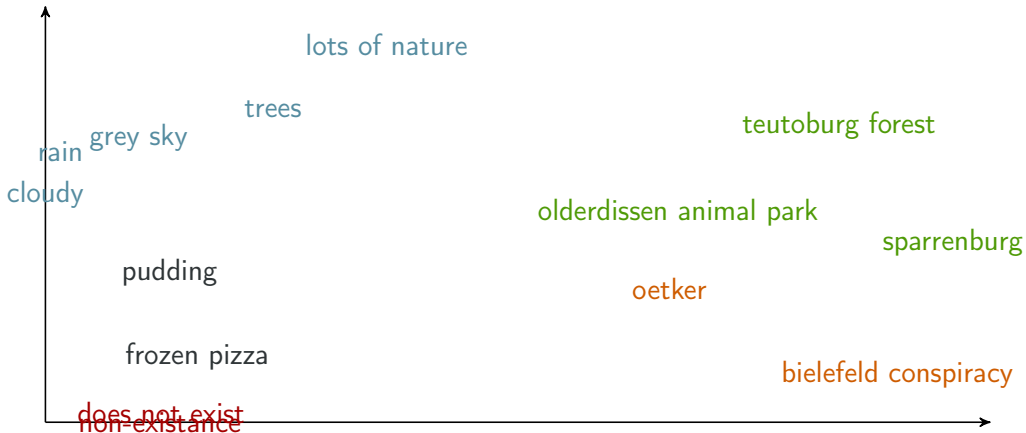
## Step 2: Clustering

- ▶ Select  $K$  via silhouette score and BIC
  - ▶ Apply relational  $K$ -Means (actually: just  $K$ -means on normalized embeddings)
- ⇒ Problem: very large number of clusters :(

## Step 3: Agglomerative clustering

- ▶ merge all clusters above similarity threshold via aggl. clustering
- ▶ re-compute cluster means
- ▶ interpret clusters via words closest to center

- PCA on word embeddings, color indicates  $K$ -means clustering with  $K = 5$



- ▶ Via aggl. clustering and relational  $K$ -means, clustering is possible even on purely distance-based data
- ▶ example: program clustering, word clustering
- ▶ challenge 1: efficiency (at least  $\mathcal{O}(N^2)$ )
- ▶ challenge 2: interpretability (e.g. via closest-to-mean, largest coefficient; Hofmann et al. 2014)

Bouguettaya, Athman et al. (2015). “Efficient agglomerative hierarchical clustering”. In: **Expert Systems with Applications** 42.5, pp. 2785–2797. doi:

10.1016/j.eswa.2014.09.054.

Cormack, R. M. (1971). “A Review of Classification”. In: **Journal of the Royal Statistical Society: Series A (General)** 134.3, pp. 321–353. doi:

10.2307/2344237.

Hammer, Barbara and Alexander Hasenfuss (2010). “Topographic Mapping of Large Dissimilarity Data Sets”. In: **Neural Computation** 22.9, pp. 2229–2284. doi:

10.1162/NECO\_a\_00012.

Hofmann, Daniela et al. (2014). “Learning interpretable kernelized prototype-based models”. In: **Neurocomputing** 141, pp. 84–96. doi:

10.1016/j.neucom.2014.03.003.

Lance, GN and WT Williams (1966). “A generalized sorting strategy for computer classifications”. In: **Nature** 212.5058, pp. 218–218. doi: 10.1038/212218a0.

Paaßen, Benjamin (2019). “Metric Learning for Structured Data”. [Dissertation](#).  
Bielefeld University. doi: [10.4119/unibi/2935545](https://doi.org/10.4119/unibi/2935545).